

Tugas UTS Machine Learning

Nama : Ilham Novriadi

NIM : 231011403539

Dataset : Bank Marketing Dataset

(<https://www.kaggle.com/datasets/janiobachmann/bank-marketing-dataset>)

Get the data

In [18]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load dataset (adjust the path if needed)
df = pd.read_csv("/Volumes/IDEKU/University/Code/data/bank_marketing.csv")

# Show basic info
print(df.shape)
df.head()
```

(11162, 17)

Out[18]:

	age	job	marital	education	default	balance	housing	loan	contact	...
0	59	admin.	married	secondary	no	2343	yes	no	unknown	
1	56	admin.	married	secondary	no	45	no	no	unknown	
2	41	technician	married	secondary	no	1270	yes	no	unknown	
3	55	services	married	secondary	no	2476	yes	no	unknown	
4	54	admin.	married	tertiary	no	184	no	no	unknown	

EDA (Exploratory Data Analysis)

In [19]:

```
# Info about data types and missing values
df.info()

# Check missing values
print("\nMissing values per column:\n", df.isnull().sum())

# Target variable distribution
print("\nTarget value counts:")
print(df['deposit'].value_counts())

sns.countplot(x='deposit', data=df)
plt.title("Deposit Target Distribution")
plt.show()

# Example of numeric feature distribution
num_cols = ['age', 'balance', 'duration', 'campaign', 'pdays', 'previous']
df[num_cols].hist(figsize=(10,8), bins=20)
```

```
plt.tight_layout()  
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 11162 entries, 0 to 11161  
Data columns (total 17 columns):  
 #   Column      Non-Null Count Dtype  
---  
 0   age         11162 non-null  int64  
 1   job          11162 non-null  object  
 2   marital      11162 non-null  object  
 3   education    11162 non-null  object  
 4   default      11162 non-null  object  
 5   balance      11162 non-null  int64  
 6   housing      11162 non-null  object  
 7   loan          11162 non-null  object  
 8   contact       11162 non-null  object  
 9   day           11162 non-null  int64  
 10  month         11162 non-null  object  
 11  duration     11162 non-null  int64  
 12  campaign     11162 non-null  int64  
 13  pdays         11162 non-null  int64  
 14  previous      11162 non-null  int64  
 15  poutcome      11162 non-null  object  
 16  deposit        11162 non-null  object  
dtypes: int64(7), object(10)  
memory usage: 1.4+ MB
```

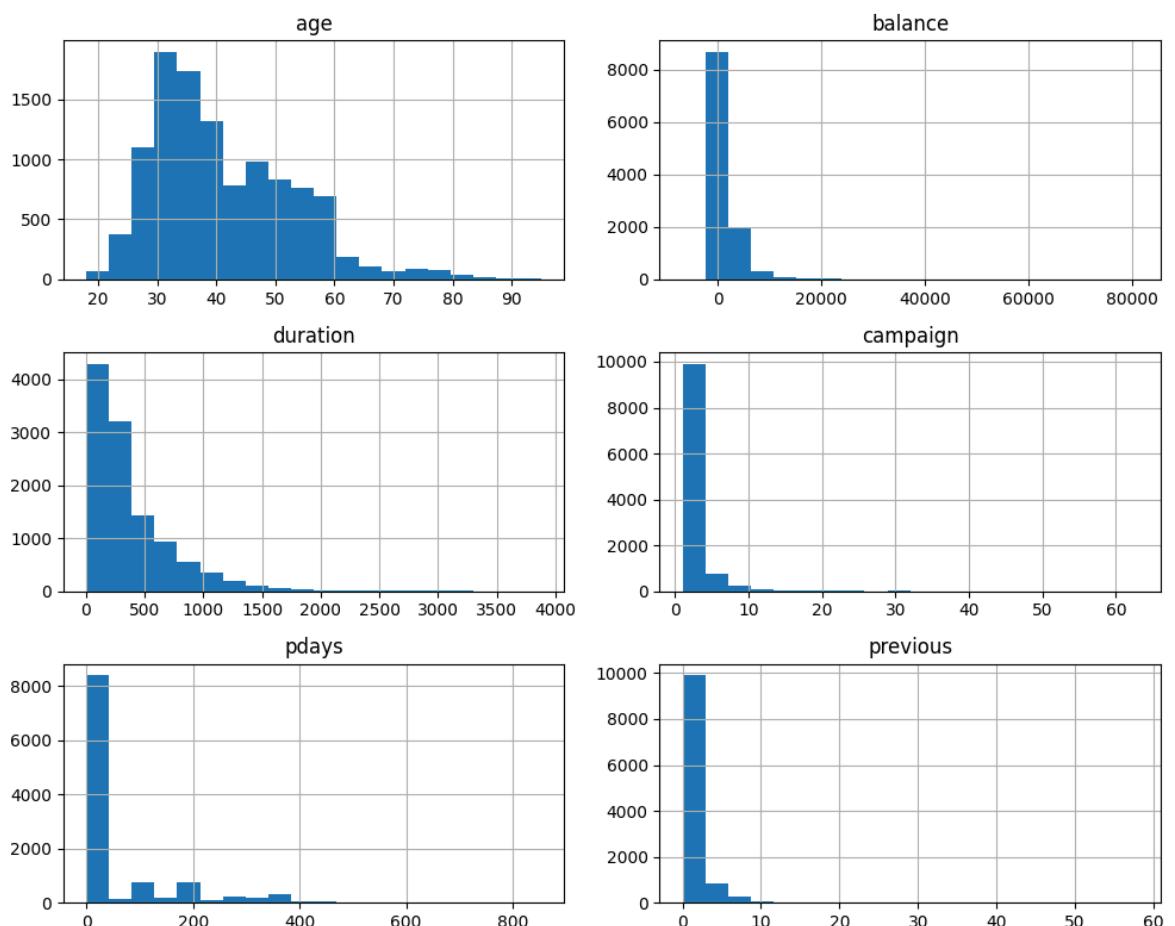
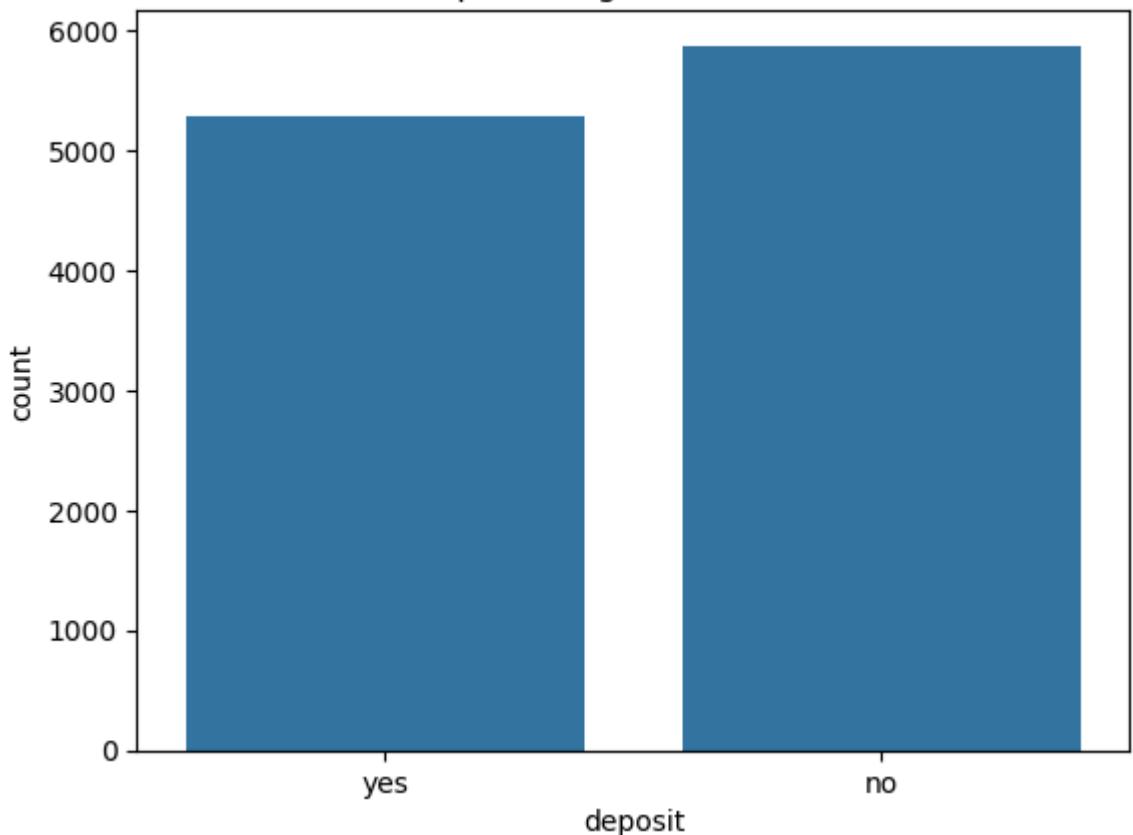
Missing values per column:

```
age          0  
job          0  
marital      0  
education    0  
default      0  
balance      0  
housing      0  
loan          0  
contact       0  
day           0  
month         0  
duration     0  
campaign     0  
pdays         0  
previous      0  
poutcome      0  
deposit        0  
dtype: int64
```

Target value counts:

```
deposit  
no      5873  
yes     5289  
Name: count, dtype: int64
```

Deposit Target Distribution



Preprocessing

```
In [20]: from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler, OneHotEncoder
```

```

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# Convert target to binary
df['deposit'] = df['deposit'].map({'yes': 1, 'no': 0})

# Split features and target
X = df.drop('deposit', axis=1)
y = df['deposit']

# Identify categorical & numeric columns
cat_cols = X.select_dtypes(include=['object']).columns.tolist()
num_cols = X.select_dtypes(include=['int64', 'float64']).columns.tolist()

# Column transformer: scale numeric, one-hot encode categorical
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), num_cols),
        ('cat', OneHotEncoder(handle_unknown='ignore'), cat_cols)
    ]
)

# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

```

Train Models

In [21]:

```

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

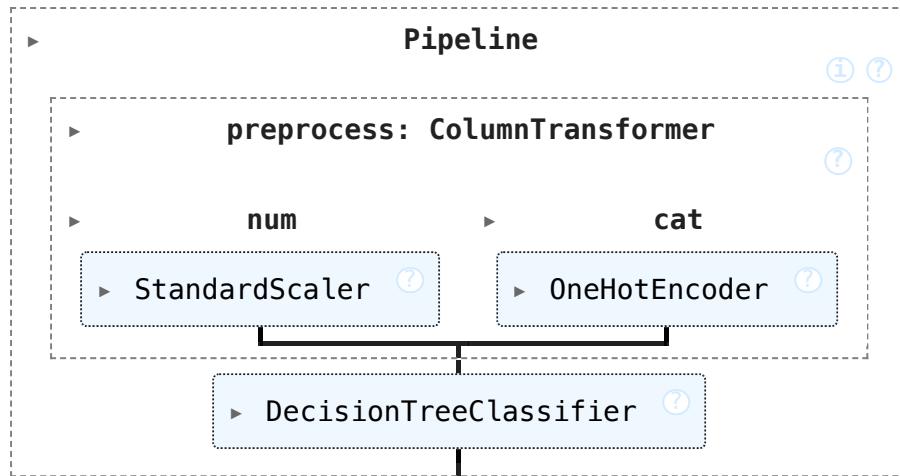
# Logistic Regression pipeline
pipe_lr = Pipeline([
    ('preprocess', preprocessor),
    ('clf', LogisticRegression(max_iter=1000, class_weight='balanced', ra
    ])

# Decision Tree pipeline
pipe_dt = Pipeline([
    ('preprocess', preprocessor),
    ('clf', DecisionTreeClassifier(max_depth=5, random_state=42, class_we
    ])

# Fit both models
pipe_lr.fit(X_train, y_train)
pipe_dt.fit(X_train, y_train)

```

Out[21]:



Evaluate Models

```
In [22]: from sklearn.metrics import classification_report, confusion_matrix, Conf

def evaluate_model(name, model):
    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:,1] if hasattr(model.named_st

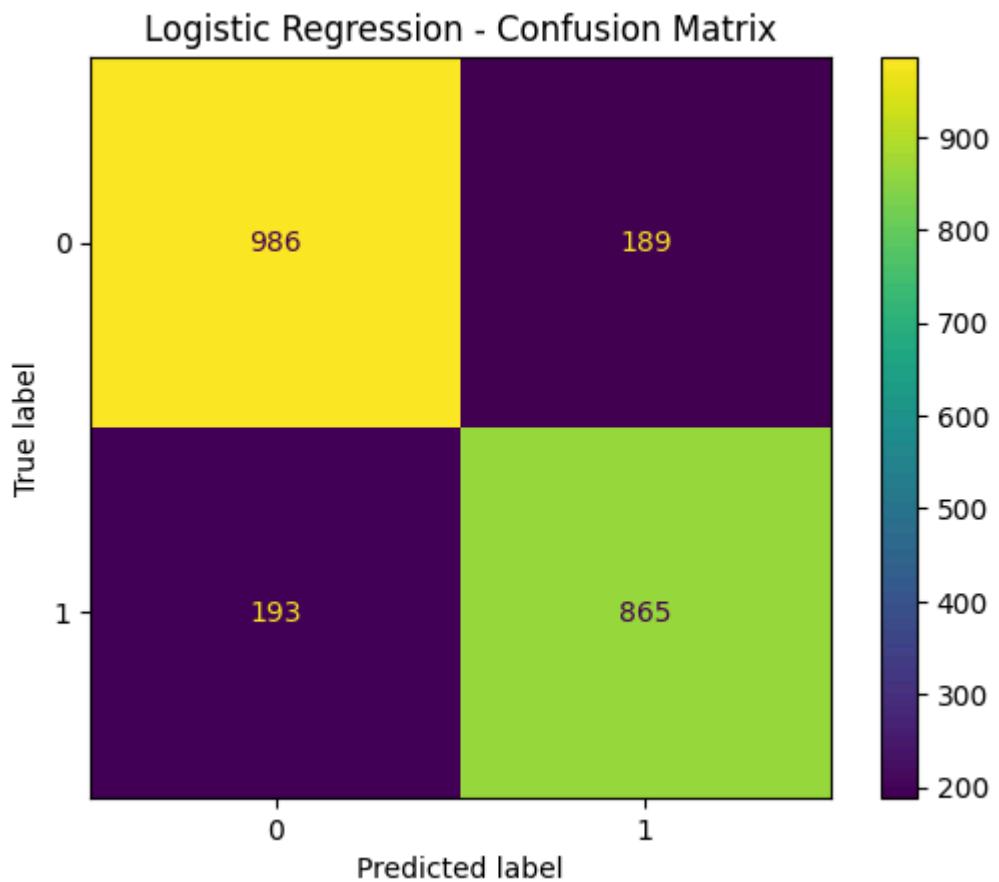
    print(f"\n== {name} ==")
    print(classification_report(y_test, y_pred, digits=4))
    ConfusionMatrixDisplay(confusion_matrix(y_test, y_pred)).plot()
    plt.title(f'{name} - Confusion Matrix')
    plt.show()

    if y_proba is not None:
        auc = roc_auc_score(y_test, y_proba)
        print(f"ROC-AUC: {auc:.4f}")
    print("="*40)

# Evaluate
evaluate_model("Logistic Regression", pipe_lr)
evaluate_model("Decision Tree", pipe_dt)

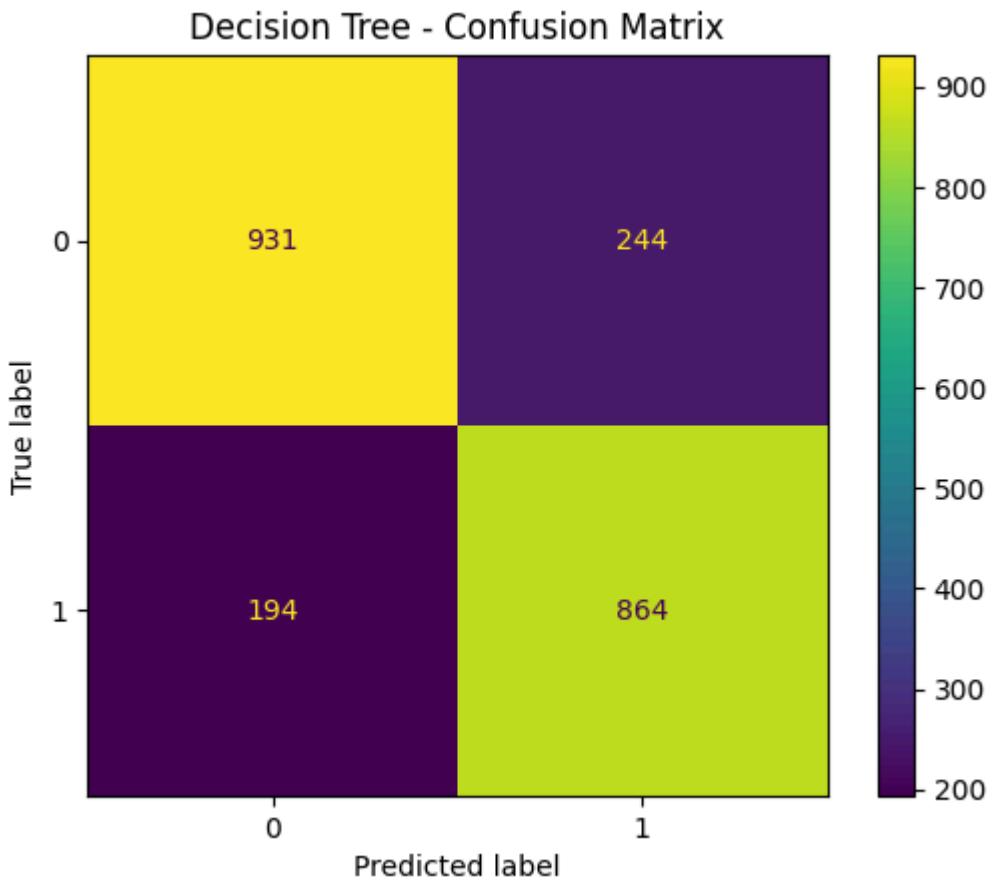
==== Logistic Regression ====
      precision    recall  f1-score   support
          0       0.8363    0.8391    0.8377     1175
          1       0.8207    0.8176    0.8191     1058

      accuracy                           0.8289     2233
     macro avg       0.8285    0.8284    0.8284     2233
  weighted avg       0.8289    0.8289    0.8289     2233
```



==== Decision Tree ===

	precision	recall	f1-score	support
0	0.8276	0.7923	0.8096	1175
1	0.7798	0.8166	0.7978	1058
accuracy			0.8039	2233
macro avg	0.8037	0.8045	0.8037	2233
weighted avg	0.8049	0.8039	0.8040	2233



ROC-AUC: 0.8745

Compare All Models

```
In [23]: models = {
    'Logistic Regression': pipe_lr,
    'Decision Tree': pipe_dt
}

results = []

for name, model in models.items():
    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:,1]
    auc = roc_auc_score(y_test, y_proba)
    f1 = np.mean([f1 for f1 in classification_report(y_test, y_pred, output_dict=True).values()])
    results.append([name, auc, f1])

results_df = pd.DataFrame(results, columns=['Model', 'ROC-AUC', 'Avg F1'])
results_df
```

	Model	ROC-AUC	Avg F1
0	Logistic Regression	0.907147	558.871686
1	Decision Tree	0.874472	558.853189

ROC Curve Visualization

```
In [24]: from sklearn.metrics import RocCurveDisplay

plt.figure(figsize=(6,5))
for name, model in models.items():
    RocCurveDisplay.from_estimator(model, X_test, y_test, name=name)
plt.title("ROC Curves Comparison")
plt.show()
```

<Figure size 600x500 with 0 Axes>

